

The table presents macro F1 scores and their arithmetic mean for selected configurations in each model family.

- Unless indicated otherwise, all models have been trained for 3 epochs.
- n_π denotes the number of random directions.
- $\text{blur} = \sqrt{\varepsilon}$, where ε is the parameter in the entropy-regularised 1-Wasserstein cost.
- Intuitively, τ penalises creating or destroying mass.
- Ensemble results: UB – Unbalanced Optimal Transport, B – Balanced OT, CL – Contrastive Learning, ep. – epoch.
- The multiplier indicates the voting weight.