

- ▶ **Blur ($\sqrt{\varepsilon}$).** Best results are achieved with $\text{blur} \approx 0.003$. Larger values lead to more bias, smaller values cause more variance, so F1 decreases in either direction.
- ▶ **Orthogonality weight.** Varying λ shifts macro F1 by $\approx < 0.003$. This implies that the OT scalar already explains the signal.
- ▶ **Style dimension.** Performance is effectively flat between 64 and 256 dimensions, which means that stylistic information is low-rank.
- ▶ **Projection dimension.** There is no gain beyond $n_\pi = 32$, and larger values simply lead to slower training.

Apart from blur tuning in the feature-free balanced OT method, OT variants are not particularly sensitive to orthogonality weight/projection dimension/style dimension tuning in the context of style shift detection.